

# ViRel: Unsupervised Visual Relations Discovery with Graph-level Analogy

Daniel Zeng\*, Tailin Wu\*, Jure Leskovec

Stanford University



## Introduction

- Visual relations form the basis of understanding our compositional world, as relationships between visual objects capture key information in a scene.
- We tackle learning and discovering relations without supervision, where relation types and labels are not known *a priori*
- In contrast, when relations are learned with predefined labels in a supervised context, this limits us to settings which depend on those seen relations

## Problem Statement

- Input:  $[16 \times 16 \times 9 \times n]$  images, and no knowledge of objects, nor relation types, nor underlying graph per example
- The goal: Infer the global relation types
- Graph  $E(g_i)$  of each image  $x_i$
- Relational graph structure  $E(G_t)$  for each task  $t$ 
  - Requires identifying relation  $e_{k,l}$  between object pairs
  - Metric: Relation Classification Accuracy %

## Dataset: BabyARC

- Collection of images (observations)  $x_i$
- Each image  $x_i$  belongs to some known task  $t \in T$ .
- Each task  $t$  has an unknown unique task graph  $G_t = (V, E)$ , where nodes  $V$  are objects and edges  $E$  are the corresponding relations
- All of its corresponding images share this common relational subgraph  $E(G_t)$
- Objects in  $x_i$  part of relational subgraph: “core” objects  
Not part of relational subgraph: “distractor” objects

Figure 1. The relations in dataset, ‘same-shape’, ‘same-color’, ‘inside’

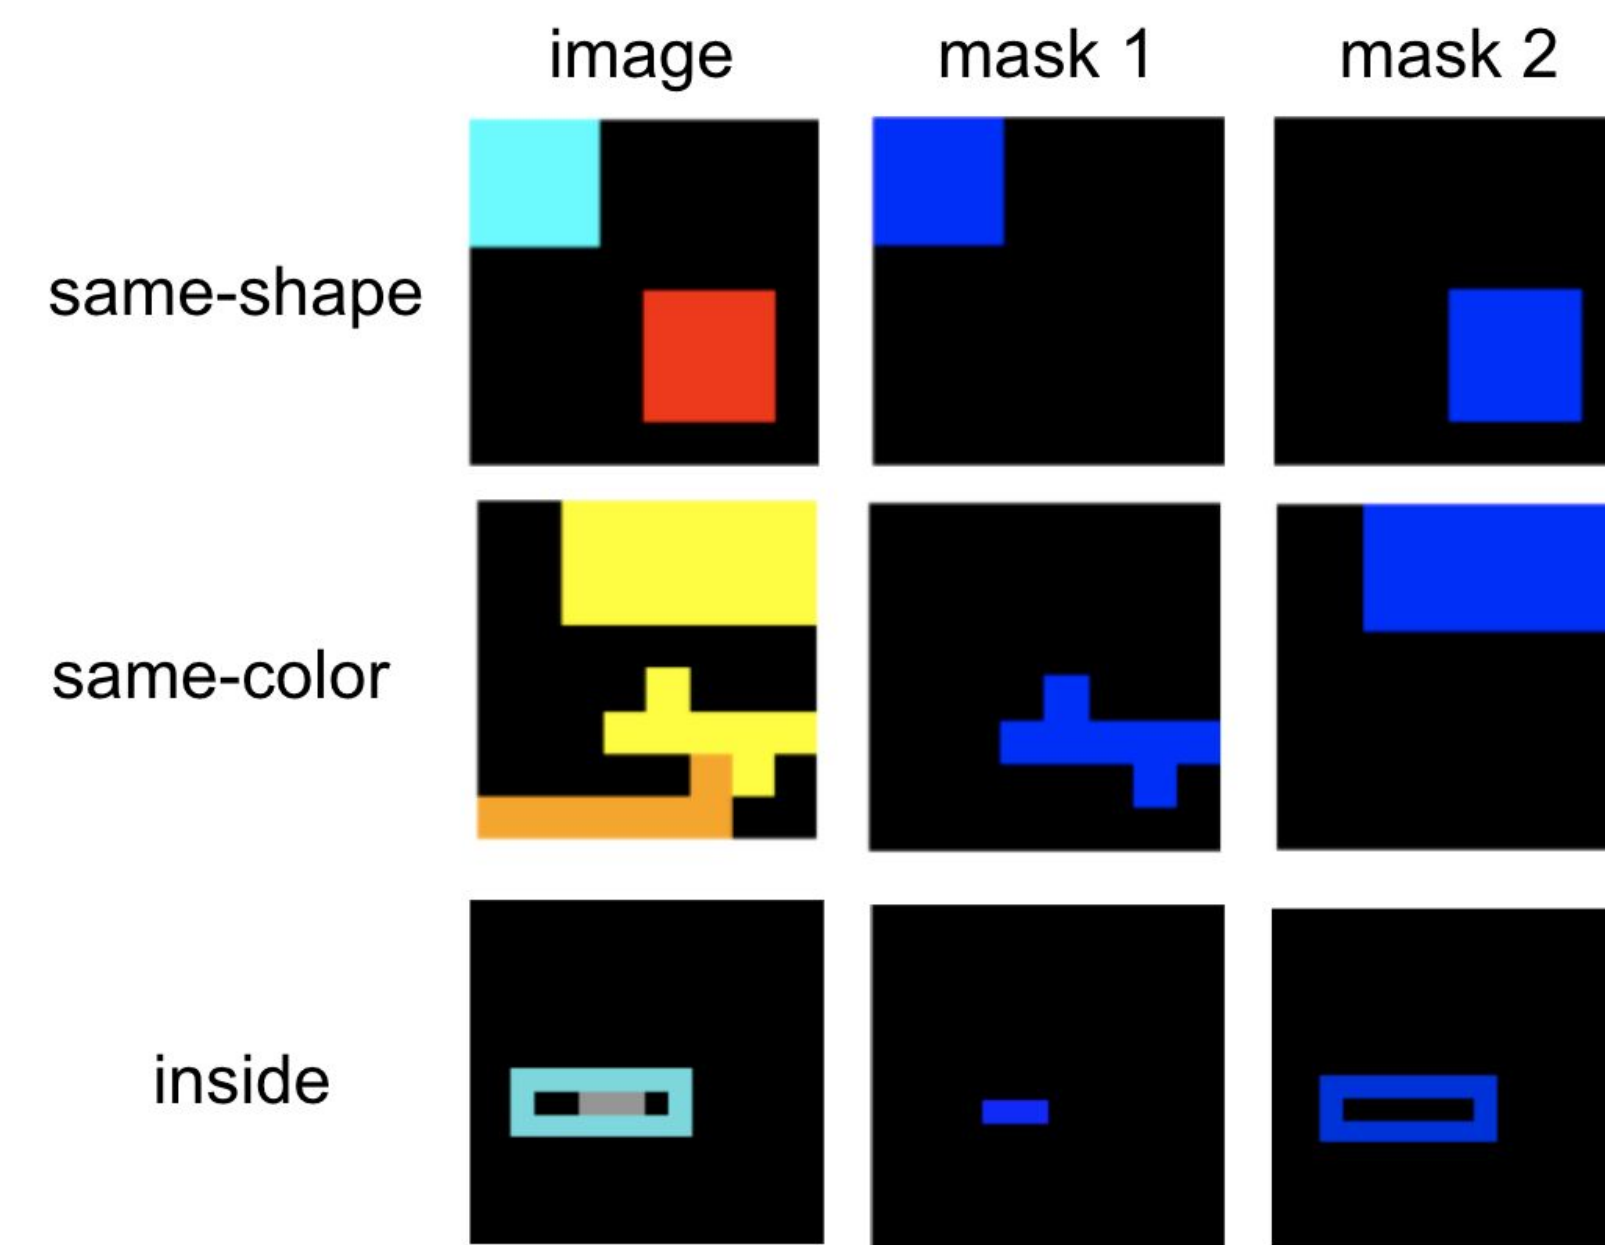

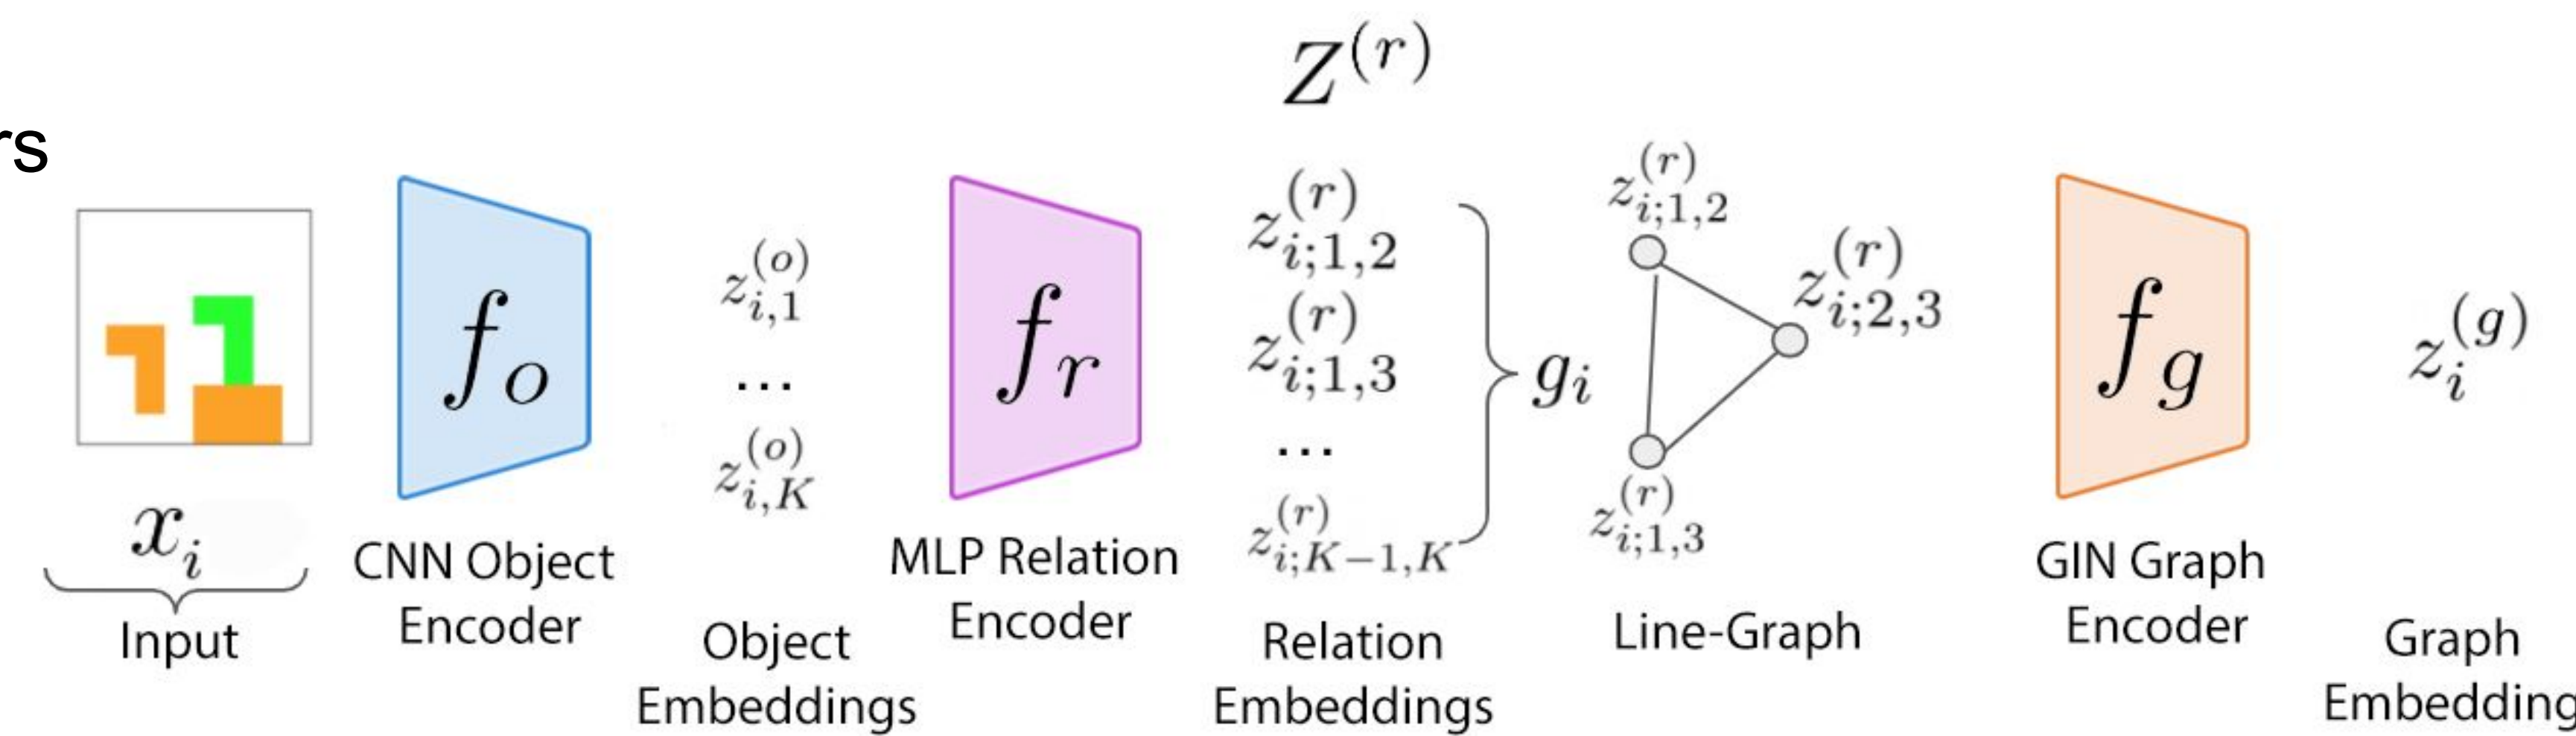


Figure 2. t-SNE visualization of relation embeddings. 0 is ‘none’, 1 is ‘inside’, 2 is ‘same-color’, 3 is ‘same-shape’



## Methods



- Each image  $x_i$  is represented as latent graph  $g_i$  after CNN + MLP, becomes  $z_i^{(g)}$  after graph isomorphism network (GIN)
- Two types of loss functions to train CR-GNN architecture (above)
  - Contrastive objective

$$\mathcal{L}_{\text{contrastive}} = \sum_{i,j \in \text{same-task}} \|f_g(g_i) - f_g(g_j)\|_2 + \sum_{k,m \in \text{diff-task}} \max(0, \eta - \|f_g(g_k) - f_g(g_m)\|_2)$$

- graph representation  $f_g(g_i)$  should be similar within the task (intra-task loss), and should be different between different tasks (inter-task loss).  $\eta$  is margin hyperparameter

- Classification objective

$$\mathcal{L}_{\text{classify}} = \sum_{\forall i \in n} \mathcal{L}_{CE}(\text{Linear}(f_g(g_i)), y)$$

- standard cross-entropy loss, between the true task ID  $y$  against the predicted task ID
- Additional regularizer: Information Bottleneck (IB)
  - Constraints information between input and relation embedding

$$\mathcal{L}_{\text{IB}} = \mathcal{I}(X; Z^{(r)})$$

## Results

Table 1. Relation classification accuracy for 2-3 core objects

METHOD	# DISTRACTORS		
	0	1	0-2
CLASSIFY	0.923	0.926	0.946
CLASSIFY + IB	0.919	0.918	0.901
CONTRASTIVE	<b>0.959</b>	0.961	0.954
CONTRASTIVE + IB	0.952	<b>0.963</b>	<b>0.957</b>
BEST	<b>0.959</b>	<b>0.963</b>	<b>0.957</b>

Table 2. Relation classification accuracy for 2-4 core objects


METHOD	# DISTRACTORS		
	0	1	0-2
CLASSIFY	0.956	0.955	0.965
CLASSIFY + IB	0.960	0.962	0.959
CONTRASTIVE	<b>0.965</b>	0.971	0.965
CONTRASTIVE + IB	0.960	<b>0.973</b>	<b>0.971</b>
BEST	<b>0.965</b>	<b>0.973</b>	<b>0.971</b>

- Apply  $k$ -means clustering to assign cluster labels to each of the learned relation embeddings
- Model performance is similar between both objectives
  - Contrastive objective performing slightly better
- No accuracy degradation due to varying the number of introduced distractor objects, in all cases
- Overall slight improvement with more tasks in 2-4 core objects setting compared to 2-3 core objects
- Model infers the global relation types, as the t-SNE visualization shows clustering of relation embeddings by same relation label, shown in Figure 2

## Conclusion / Future Work

- Our method achieves above 95% accuracy in relation classification, discovers the relation graph structure for most tasks, and further generalizes to unseen tasks with more complicated relational structures
- Limitation: model only learns the necessary relation representations needed to distinguish between the given tasks, such as overlapping clusters in Figure 2
- Future work is expanding towards more datasets, potentially CLEVR generation with graph structure

### Example Tasks



### Task Graph $E(G_t)$

